


Sentence Embedding Similarity: Domain Analysis (BGE-large model, consecutive sentence pairs)

Embedding Variance Per Domain
(higher = more discriminative boundary signal)



Distribution of Consecutive Similarities
(MATH flat = low boundary discriminability)

